



Article

VENTINEL: Automated Detection of Android Vishing Apps Using Optical Character Recognition

Daegyeom Kim ¹, Sehwan O ¹, Younghoon Ban ¹, Jungsoo Park ², Kyungho Joo ¹ and Haehyun Cho ^{1,*}

¹ Graduate School of Software, Soongsil University, Seoul 06978, Republic of Korea; daegyeom88@soongsil.ac.kr (D.K.)

² Department of ICT Convergence Science, Kangnam University, Yongin 16979, Republic of Korea

* Correspondence: haehyun@ssu.ac.kr; Tel.: +82-2-828-7360

Abstract: Vishing, a blend of “voice” and “phishing”, has evolved to include techniques like Call Redirection and Display Overlay Attacks, causing significant financial losses. Existing research has largely focused on user behavior and awareness, leaving gaps in addressing attacks originating from vishing applications. In this work, we present VENTINEL, an Android-based defense system designed to detect these attacks without requiring OS modifications. VENTINEL employs Optical Character Recognition (OCR) to compare phone numbers during calls, effectively preventing Call Redirection and Display Overlay Attacks. Additionally, it safeguards against Duplicated Contacts Attacks by cross-referencing call logs and SMS records. VENTINEL achieves 100% detection accuracy, surpassing commercial applications, and operates with minimal data collection to ensure user privacy. We also describe malicious API behavior and demonstrate that the same behavior is possible for API levels 29 and higher. Furthermore, we analyze the limitations of existing solutions and propose new attack and defense strategies.

Keywords: Android; mobile security; voice phishing; vishing



Academic Editors: Weizhi Meng and Christian D. Jensen

Received: 25 November 2024

Revised: 2 January 2025

Accepted: 4 January 2025

Published: 7 January 2025

Citation: Kim, D.; O, S.; Ban, Y.; Park, J.; Joo, K.; Cho, H. VENTINEL: Automated Detection of Android Vishing Apps Using Optical Character Recognition. *Future Internet* **2025**, *17*, 24. <https://doi.org/10.3390/fi17010024>

Copyright: © 2025 by the authors. Licensee MDPI, Basel, Switzerland. This article is an open access article distributed under the terms and conditions of the Creative Commons Attribution (CC BY) license (<https://creativecommons.org/licenses/by/4.0/>).

1. Introduction

Vishing, a portmanteau of “Voice” and “Phishing”, has evolved with advancements in technology, allowing attackers to employ additional phishing channels such as web, Short Message Service (SMS), and email for more sophisticated scams [1]. These techniques have continued to result in annual losses amounting to millions of dollars for both individuals and organizations [2,3]. The emergence of attacks utilizing vishing applications has significantly escalated the fraud landscape; leveraging these applications has been shown to incur losses up to ten times greater than those caused by traditional phishing campaigns that do not utilize such apps [4].

With the proliferation of vishing applications, the most commonly observed attack methods have been the Call Redirection Attack and the Display Overlay Attack [4]. Firstly, the Call Redirection Attack involves redirecting the number dialed by the user to the attacker’s phone number. However, since the attacker’s phone number would appear on the user’s outgoing call screen, this method is often combined with a Display Overlay Attack for added deception. The Display Overlay Attack, on the other hand, is used to conceal the attacker’s phone number by overlaying a fake waiting screen during incoming or outgoing calls [5]. By masking the real phone number with a counterfeit screen, users are less likely to notice discrepancies, thereby increasing the likelihood of them divulging personal information without suspicion. Additionally, we defined a new attack: Duplicated Contacts Attack. This attack involves registering the attacker’s phone number under a

contact name that matches a target contact's name or editing an existing contact on the target's device to insert the attacker's phone number.

In previous research, due to the lack of clearly defined vishing attack techniques, researchers explored various dimensions of vishing and investigated user behavior in response to these threats. Based on these studies, research has been conducted to enhance users' security awareness, focusing on enabling individuals to recognize and respond to phishing attempts. These include user studies analyzing why individuals respond (or do not respond) to vishing attempts [6], research aimed at increasing awareness of mobile scams through education [7], and detailed crime reports explaining how attackers succeeded in their fraudulent schemes [8]. However, the emergence of vishing applications has introduced new attack vectors, such as Call Redirection and Display Overlay Attacks, leaving gaps in the existing body of research. To fill these gaps, an approach involving modifications to the Android Operating System (OS) was proposed, which monitors specific Application Programming Interfaces (API) to detect malicious behaviors [4]. However, this method requires users to directly alter the Android OS, which poses a significant barrier to adoption. Additionally, it is limited to monitoring specific APIs used in malicious techniques. If Google introduces new APIs with similar functionalities, this approach would fail to detect them. There has also been research utilizing Large Language Models (LLMs) to assist victims in responding to vishing attempts. However, these approaches face challenges due to limited training data, which introduces biases and reduces the reliability of the results [9].

In this work, VENTINEL focuses on monitoring the user's screen to defend against Display Overlay and Call Redirection Attacks. When an incoming call is detected, VENTINEL uses Optical Character Recognition (OCR) to compare the number displayed on the user's screen with the incoming number to detect the Display Overlay Attack. Additionally, when an outgoing call is initiated, VENTINEL compares the number dialed by the user with the number that successfully connects, which allows it to detect the Call Redirection Attack. Then, through OCR, VENTINEL checks for the Display Overlay Attack that could be used alongside Call Redirection Attack by comparing the number displayed on the screen with the number dialed by the user. Finally, VENTINEL examines the user's contacts and cross-references the phone numbers stored in the recent call log and SMS records to identify unused contacts, thereby defending against Duplicated Contacts Attacks. VENTINEL stores the identified phone numbers in the blacklist and, after confirming the user's intention to delete them, proceeds with deleting the contacts. In the process, VENTINEL does not collect or store unnecessary data and only accesses the minimum data needed to monitor duplicate contacts by asking for explicit permission from the user.

This method emphasizes pre-incident detection and operates entirely within the Android application, ensuring user convenience. Furthermore, even if the APIs used to execute the same malicious actions are changed, the VENTINEL remains capable of detection. VENTINEL demonstrated a 100% detection rate for Display Overlay, Call Redirection, and Duplicated Contacts, achieving higher scores compared to commercial apps available on the Android Play Store. Additionally, we conducted a user study using a benchmark that we developed, which revealed that, when using VENTINEL, only approximately 8.9% of users responded on average.

We present the following contributions:

- We analyze APIs provided in Android level 29 and above that exhibit behaviors similar to malicious activities.
- We identify and highlight the limitations of currently available commercial vishing defense applications.

- We examine existing attack techniques used by vishing applications and propose novel attack methods.
- We implement a defense mechanism within an application to counteract vishing attacks, specifically focusing on Display Overlay Attacks, Call Redirection Attacks, and the newly defined Duplicated Contacts Attack, all without requiring OS modifications.

2. Related Work

Various solutions exist to mitigate the impact of vishing, which are categorized into pre-incident and post-incident techniques. This section provides an overview of vishing defense methods, dividing them into pre-incident and post-incident approaches.

2.1. Pre-Incident Detection

2.1.1. Typical Approaches

To prevent vishing and mobile fraud, governments and corporations typically provide mobile security training [7]. Mobile security education aims to enhance cybersecurity awareness among general users, helping them become more familiar with security terminology. However, Chin et al. [10] argue that there is little difference in security awareness between groups that received mobile security training and those that did not, suggesting that such training alone is not effective in blocking vishing.

2.1.2. API-Trace-Based Approaches

HearMeOut [4] was proposed as a modified version of AOSP to automatically detect and alert users about otherwise unnoticed vishing attempts. HearMeOut monitors three features implemented using the functions listed in Table 1, which are utilized by vishing applications. Through this, HearMeOut detects Call Redirection and Call Screen Overlay as pre-incident detection, and detects Fake Call Voice as a post-incident detection.

Table 1. Specific functions monitored by HearMeOut.

	broadcastIntent ()
CallRedirection	setResultData ()
	onCreateOutgoingConnection ()
	WindowManagerImpl.addView ()
Call Screen Overlay	onCreateOutgoingConnection ()
	onCreateIncomingConnection ()
	onCallStateChanged ()
Fake Call Voice	onAudioFocusChanged ()
	onPlaybackStateChanged ()

HearMeOut recruited a total of 45 Android users for evaluation. Among them, 23 participants did not use HearMeOut, while 22 participants used it. The experiment was conducted by having participants make calls to the official Woori Bank service center using pre-installed Pixel 2 phones. Ultimately, HearMeOut achieved a 100% detection rate for the three techniques and recorded zero false positives. However, it has limitations, as it is vulnerable to changes, deletions, or additions to the APIs and functions supported by Android. Addressing this issue requires support from AOSP maintainers and collaboration with Android.

2.2. Post-Incident Detection

2.2.1. Voice-Recognition-Based Approaches

Viking [9] was proposed a method designed to assist victims by leveraging AI to recognize and respond to real-time conversations during incoming and outgoing calls (i.e., Social Engineering Attacks). The operational process of Viking involves collecting information through Text-To-Speech (TTS) when a call occurs and relaying this information to a trained LLM for risk assessment of the conversation. For instance, if language patterns indicative of sensitive information requests or urgent actions are detected, the conversation is deemed risky, and the system alerts the victim, providing them with suggested responses. Ultimately, Viking aims to help victims protect their privacy and safely terminate conversations with attackers. However, Viking faces limitations due to the restricted dataset used for training the LLMs. Additionally, utilizing an LLM demands substantial resources, posing significant challenges for real-time processing and potentially leading to delays in response generation.

2.2.2. Blacklist-Based Approaches

The applications that failed to detect benchmarks in Section 5.3 all operate on a blacklist basis. Blacklist-based approaches are easy to implement and have the advantage of being able to block any phone numbers used for vishing once they are registered. However, until the numbers used by attackers are included in the blacklist, victims continue to incur losses. Moreover, attackers easily create new phone numbers using burner phones, making it challenging to mitigate the impact of vishing attacks.

2.3. Current Vishing Defenses

In practice, attackers circumvent safeguards by sharing vishing applications through websites or emails, rather than through application marketplaces like the Google Play Store [7]. Due to the difficulty in preventing such distribution, companies offer applications that provide blacklist-based warnings or utilize Android APIs to present phone numbers. The Google Play Store hosts vishing defense applications, including WhoWho [11], Truecaller [12], SafeVoice [13], Hiya [14], Whoscall [15], Phishingeyes [16], CallApp [17], Show Caller ID & Spam Blocker [18], and SmartAntiPhishing [19]. Blacklist-based defenses add the attacker's number to the blacklist after the victim has already suffered an incident, while defense applications utilizing Android APIs simply verify Call Redirection based on the numbers. These found in Table 2. However, these defense techniques are not effective against new vishing applications or Display Overlay Attack.

Table 2. Vishing defense methods advantages and disadvantages.

Blacklist-based approaches	
Advantages	100% detection when phone number exists in list
Disadvantages	Damage caused by post incident detection
Voice recognition-based approaches	
Advantages	Detecting voice phishing through specific text
Disadvantages	LLM dataset limitations and response generation delays
Provide number-based approaches	
Advantages	Enable Call Redirection Attack defense
Disadvantages	Warning is passive and there is no direct defense
API trace-based approaches	
Adavantages	Various API call information checked
Disadvantages	Impossible to detect when API changes

3. Background

In this section, we introduce the malicious behaviors found in vishing applications and the Android APIs that enable the creation of these behaviors. We also explain how vishing applications use Android APIs differently depending on the Android API level. Furthermore, we discuss known vishing defense approaches and their limitations.

3.1. Malicious Behavior in Vishing Apps

Vishing, a combination of “Phishing” and “Voice”, refers to a type of mobile cyber crime that uses voice calls to financially exploit victims [20,21]. Attackers lure victims into installing vishing applications and subsequently impersonate public institutions or acquaintances through these applications. Consequently, victims are misled by insidious fraud and inadvertently provide sensitive information, potentially leading to financial harm [22]. The malicious behaviors identified in the vishing applications we examined are as follows:

3.1.1. Call Redirection

Android provides APIs for Call Redirection, enabling developers to create applications that cancel outgoing calls and automatically connect to a different phone number. Therefore, attackers develop vishing applications to cancel a victim’s requested outgoing call and redirect it to an attacker. To implement Call Redirection, attackers employ different APIs based on the target Android API level. We detail different Android APIs corresponding to API levels in Section 3.2. Additionally, vishing applications combine Display Overlay with Call Redirection to disguise the redirected outgoing number.

3.1.2. Display Overlay

Android supports APIs that facilitate the output of specific pop-up screens or full-screen displays, known as Display Overlay [23,24]. Vishing applications exploit these APIs to perform malicious activities by using Display Overlay to cover the victim’s device screen with a fake interface, thereby deceiving the victim about the outgoing caller ID. Specifically, when vishing applications change the outgoing call, they overlay a fake screen to disguise the altered outgoing number display. When victims receive a call from attackers, vishing applications overlay a screen that appears to be from a legitimate institution (e.g., police, bank, public institution) to deceive the victim.

3.1.3. Duplicated Contacts

Android provides APIs that allow developers to retrieve, and modify contacts stored on a device. Vishing applications exploit these APIs to exfiltrate or modify contacts on a victim’s device for malicious purposes. Additionally, we introduce a malicious behavior called *Duplicated Contacts*. Duplicated Contacts involves adding the attacker’s number to the victim’s contact list under an existing contact name or editing it as a secondary contact. According to a FINLEY, Jason R., et al. [25], many smartphone users store contacts on their devices rather than memorizing phone numbers, making Duplicated Contacts effective in tricking victims into thinking they are receiving a call from someone they know. Android users overlook the actual number displayed beneath the contact name, and iOS users do not see the number at all, allowing attackers to exploit this limitation.

3.1.4. Application Repackaging

Application repackaging is a technique that involves decompiling an application, modifying its code, and recompiling it into a distributable format (such as Android Application Package (APK) or iOS Application Archive (IPA)). Attackers exploit this method by decompiling signed applications, adding malicious code, and repackaging them [26].

The repackaged applications maintain the layout of the original applications, making it difficult for victims to detect the alterations [27]. Attackers have deceived victims by utilizing vishing applications that encompass the previously mentioned behaviors. Experts caution that individuals and businesses incur annual financial losses amounting to millions of dollars as a result of vishing-related fraud [2,3]. To prevent vishing fraud, researchers have proposed various approaches to detect vishing behaviors using different methods [4,28–30]. Unfortunately, despite these efforts, incidents of financial loss resulting from vishing continue to rise [31]. Therefore, we propose an approach to detect malicious behaviors, specifically Display Overlay, Call Redirection, and Duplicated Contacts, without modifying the Android OS.

3.2. Vishing Apps According to Android API Levels

The Android platform provides various APIs (e.g., TelephonyManager API, LocationManager API, etc.) that request information from the device or manipulate it directly, facilitating flexible feature development for applications. Attackers exploit these APIs to create malicious applications, such as vishing applications. Google has recognized these issues and has been consistently implementing security patches [32,33], as well as removing APIs that were exploited by malicious apps, restricting the values that apps request from APIs, and requiring additional permission checks before app installation, among others. Specifically, Google introduced security patches in the API level 29 update to prevent the abuse of Call Redirection [34]. In this section, we introduce the alterations in Call Redirection and Display Overlay Attacks in vishing applications based on this security patch.

Lower API Level 29. In Android API levels 28 and below, vishing applications implement Call Redirection using `BroadcastReceiver`. At this level, the Android system broadcasts `NEW_OUTGOING_CALL` just before an outgoing call is made. Consequently, vishing applications retrieve or modify the broadcast value in the `onReceive()` callback function, thus invoking `getResultSetData()` to obtain the victim's phone number and using `setResultSetData()` to redirect the call. Vishing applications involve Display Overlay to present a counterfeit screen when displaying specific call numbers. Therefore, these applications employ APIs to detect phone events and retrieve the incoming or outgoing numbers. Vishing applications utilize the `onCallStateChanged()` and `onCallAdded()` callback functions to retrieve call states and incoming number, as shown in Table 3. Specifically, `onCallStateChanged()` provides both incoming and outgoing phone numbers through `EXTRA_INCOMING_NUMBER`.

Table 3. Constants that change depending on the call state in Android.

Call State	Constant
Incoming Call	<code>TelephonyManager.EXTRA_STATE_RINGING</code>
Outgoing Call	<code>TelephonyManager.EXTRA_STATE_OFFHOOK</code>
Waiting Call	<code>TelephonyManager.EXTRA_STATE_IDLE</code>

API Level 29 and Above. In Android API level 29 and above, Android modified the call-related APIs and categorized the phone permissions in more detail. First, to prevent Call Redirection that altered broadcast values, Android restricted access to broadcast values and introduced the `CallRedirectionService`. Android modified the behavior of the `BroadcastReceiver` so that the `setResultSetData()` method called just before an outgoing call returned `null`. Furthermore, for an application to utilize the `CallRedirectionService`, user consent is required, and the application must be registered as the default *call redirecting app*. Second, Android restricted the APIs

that provided incoming or outgoing numbers in API level 28 and below by introducing the `CallScreeningService`. In earlier versions, Android apps with permissions related to `TelephonyManager` could call APIs within callback functions to obtain call status as well as incoming or outgoing numbers. In these versions, Android added the `CallScreeningService` to manage call events and introduced settings for the default app for *caller ID and spam*. Therefore, for an Android app to access the call numbers, it must implement the `CallScreeningService` and obtain user consent to be registered as the default app. Finally, the APIs required to execute Display Overlay are available across all Android API levels.

In this section, we present an approach for detecting malicious behavior in vishing applications running in the background. Our approach identifies malicious behaviors, including Call Redirection, Display Overlay, and Duplicated Contacts.

3.3. Call Redirection Detection

Vishing applications include Call Redirection Attacks to alter the outgoing call on the victim's phone using `setResultData()` or `CallRedirectionService`. These applications also implement Display Overlay Attacks that obscure changes to the outgoing number, making it challenging for victims to detect Call Redirection. To detect Call Redirection Attacks, we retrieve and compare both the original outgoing number and the redirecting outgoing number using Android APIs. Additionally, we obtain the number displayed on the call standby screen and compare it with the actual outgoing number for double verification. This approach reduces the false positives, thereby protecting users from both inadvertent and malicious call alterations.

3.4. Display Overlay Detection

Vishing applications employ Display Overlay Attacks that present a fake screen during call events (incoming or outgoing), preventing victims from identifying the actual numbers. To detect these attacks, we compare the number displayed on the screen with the actual number retrieved using Android APIs. We capture a screen during call events and extract the phone number from the captured image using OCR. We also retrieve the actual incoming or outgoing number by invoking the Android APIs in the background application, and then we compare the extracted number from the screen with the actual number. Our approach enhances the detection of Call Redirection Attacks by double verification.

3.5. Duplicated Contacts Detection

Vishing applications conduct Duplicated Contact Attacks to make victims mistake an attacker's number for an existing contact. Adding the attacker's contact using a name that matches the target or editing the target's contact to include the attacker's phone number (i.e., the secondary phone number) seen in Figure 1. These applications generate Duplicated Contacts by adding the attacker's number either under an existing contact name or as a secondary phone number to an existing contact. To counter this, we identify contacts with duplicate names or secondary phone numbers and prompts the user to review and delete any suspicious duplicates.

To effectively counter Call Redirection and Display Overlay Attacks, we searched for real phone numbers and distinguished them from those obtained through OCR. We also adjusted the timing of number comparisons to account for network delays, as OCR begins when incoming and outgoing call events are triggered, and the actual phone number is retrieved.




Figure 1. (Left) shows adding the attacker's phone number with the same name as the target. (Right) shows editing the attacker's phone number in the target contact.

4. Design

In this section, we present VENTINEL, a system designed to detect the malicious behavior of vishing applications without necessitating modifications to the Android OS. As detailed in Figure 2, VENTINEL operates as a background application, detecting malicious activities from vishing applications in real time and promptly notifying the user. VENTINEL is composed of two core modules: SENTRY and SIGNALLER. SENTRY is continuously monitors vishing applications for signs of malicious behaviors, and SIGNALLER sends alerts to users when SENTRY transmits information.




Figure 2. VENTINEL: Vishing behavior detection process on background.

4.1. SENTRY of VENTINEL

The SENTRY of VENTINEL operates in the background to monitor the malicious behavior of vishing applications. It utilizes Android APIs to gather various types of information, such as incoming and outgoing numbers, to verify the presence of malicious activities.

SENTRY detects specific malicious behaviors, including Call Redirection, Display Overlay manipulation, and Duplicated Contacts. When the Android device is on IDLE and not engaged in a call, SENTRY performs a Duplicate Contacts Verification. Upon the user's attempt to initiate a call, SENTRY conducts a Call Redirection Verification. Following this, SENTRY carries out a secondary check through Display Overlay analysis. It captures the screen displayed by the device during both incoming and outgoing call states and performs a Display Overlay Verification to detect any potential malicious activity.

4.1.1. Duplicated Contacts Verification

The SENTRY module conducts a search for duplicate contacts within the contacts stored on the Android device while in IDLE. Any contacts identified as duplicates are forwarded to the SIGNALLER module for further action. Vishing applications add or modify contacts without the victim's awareness, causing the attacker's number to appear on the incoming or outgoing call screen.

Algorithm 1 is pseudocode that enables the Sentry of VENTINEL to perform Duplicate Contact Verification. The SENTRY accesses the complete contact list, SMS messages, and call logs from the Android device. To facilitate SENTRY, the following permissions must be granted by the user explicitly:

- READ_CONTACTS
- READ_SMS
- READ_CALL_LOG

Algorithm 1 Duplicated Contact Verification.

```

1: contacts ← List of existing Contacts
2: contactMap ← initialize empty hash table
3: for contact in contacts do
4:   name ← name of contact
5:   pNumList ← phone numbers of contact
6:   contactMap[name] ← pNumList
7:   if already name in key then
8:     Add contactMap[name] to pNumList
9:   end if
10: end for
11: dupMap ← initialize empty hash table
12: for each key, value in contactMap do
13:   if value.length > 1 then
14:     dupMap[key] ← value
15:   end if
16: end for
17: historyList ← List of Recent Call and SMS history
18: whiteList ← empty set
19: for history in historyList do
20:   pNum ← phone number in history
21:   Add pNum to whiteList
22: end for
23: for each key, value in dpMap do
24:   for number in value do
25:     if number ∈ whiteList then
26:       Delete value in dpMap
27:     end if
28:   end for
29: end for

```

Once these permissions are obtained, the SENTRY identifies any contacts with the same name or multiple phone numbers associated with a single name as duplicates. Additionally, phone numbers that exist in recent call and text records are created as a whiteList, assuming that they are related to the user. The SENTRY then removes these whitelisted numbers from the identified duplicates, generating a list of suspicious duplicate contacts, which is subsequently forwarded to the SIGNALLER.

4.1.2. Call Redirection Verification

The SENTRY is designed to detect Call Redirection Attacks that intercept calls initiated by vishing applications and to alert the victim accordingly. To accomplish this, it utilizes Android APIs to request and validate call-related information from the device initiating the call. As such, SENTRY requires the following permissions:

- READ_PHONE_STATE
- READ_CALL_LOG

For devices operating at API level 29 and lower, an additional permission, PROCESS_OUTGOING_CALLS, is necessary, while devices operating at API level 29 and above require the permission READ_PHONE_NUMBER. When a phone number is entered and the dial button is pressed in Android dialing applications, the Operating System triggers the ACTION_NEW_OUTGOING_CALL action and transmits EXTRA_PHONE_NUMBER, which contains the dialed number as an intent value. The SENTRY requests this value and stores it in a designated variable for comparison against the phone number obtained when the call state transitions to EXTRA_STATE_OFFHOOK, as this value remains unaffected by any Call Redirection Attack. Subsequently, when the call state changes to EXTRA_STATE_OFFHOOK, the intent value is cleared, and the final dialed phone number is retrieved using getResultData(). The initially stored value (the intended phone number) is then compared with the final phone number received after dialing. The result of this comparison is forwarded to the SIGNALLER. This approach is effective for detecting Call Redirection Attacks across both API levels 28 and below as well as 29 and above.

4.1.3. Display Overlay Verification

Algorithm 2 shows the pseudocode for VENTINEL's SENTRY to conduct Display Overlay Verification. To monitor for Display Overlay, SENTRY compares the phone numbers displayed on the screen during incoming and outgoing calls with the actual phone numbers being dialed. To initiate this process, SENTRY first captures the screen at the beginning of an incoming or outgoing call and utilizes ML Kit [35], an OCR tool provided by Google, to extract the phone number displayed on the screen. This extraction process reads text that conforms to the phone number format from the top left to the right of the screen. Simultaneously, SENTRY employs Android APIs to retrieve the outgoing or incoming number that the device is attempting to dial. For outgoing calls, it retrieves the number from the ACTION_NEW_OUTGOING_CALL action, while for incoming calls, it uses getStringExtra(incomingNumber). Finally, SENTRY compares the two phone numbers obtained during the call event to determine whether a Display Overlay has occurred and forwards the results to the SIGNALLER. To facilitate screen capture, Ventinel sets the foreground service type to mediaProjection and requires the following permissions:

- READ_PHONE_STATE
- READ_CALL_LOG

Algorithm 2 Display Overlay Verification of incoming call.

```

1: bitmap ← Captured Screenshot
2: recogNum ← Extracted Phone Number from bitmap
3:
4: state ← Phone Call State
5: if state is EXTRA_STATE_RINGING then
6:   inNum ← Incoming Phone Number
7:   if inNum is not recogNum then
8:     Detect Malicious Behavior
9:   end if
10:  end if
11:  if state is EXTRA_STATE_OFFHOOK then
12:    outNum ← Outgoing Phone Number
13:    if outNum is not recogNum then
14:      Detect Malicious Behavior
15:    end if
16:  end if

```

For devices operating at API level 29 and above, an additional permission, READ_PHONE_NUMBER, is necessary. To implement the APIs that verify the outgoing number, the permissions detailed in Section 4.1.2 must be obtained. Moreover, to ensure VENTINEL can run in the background together with other applications, it necessitates the FOREGROUND_SERVICE permission.

4.2. SIGNALLER of VENTINEL

In VENTINEL, the SIGNALLER is responsible for notifying users of the malicious behaviors identified by the SENTRY in real time. The SIGNALLER generates a vibration on the device and displays a warning pop-up to alert the user of any detected Display Overlay or Call Redirection Attacks. To address Duplicated Contacts Attacks, the SIGNALLER presents the user with a list of identified suspicious contacts and requests their consent to delete these entries. Upon receiving the user's approval, the SIGNALLER proceeds with the deletion of the specified contacts.

5. Evaluation

In this section, we evaluate the overhead of VENTINEL and present benchmarks to compare its detection capabilities against those of commercial applications. Furthermore, we perform a manual analysis of selected commercial applications to identify key factors affecting their detection rates. To gauge VENTINEL's reliability from a user perspective, we also conducted a user study with 200 Android users.

5.1. Evaluating Overhead

We evaluated the overhead to understand how VENTINEL impacts the user's device when activated. Using a Galaxy Jean device, we compared CPU and memory usage with and without VENTINEL running. When VENTINEL was not running, CPU utilization averaged 14%, and memory usage was 74.2%. With VENTINEL running, CPU utilization increased slightly to 15%, while memory usage remained unchanged at 74.2%. These results indicate that VENTINEL does not cause any noticeable performance degradation on the device. Further details are provided in Figure 3.




Figure 3. CPU and memory usage with and without VENTINEL.

5.2. Benchmark with Vishing Malicious Behavior

To evaluate VENTINEL, we require vishing applications; however, we opt to use a benchmark with malicious behaviors, given the typically short lifecycle of vishing applications. Consequently, we develop a benchmark application at API level 28 and below and API level 29 and above that incorporates the following malicious behaviors: Display Overlay, Call Redirection, and Duplicated Contacts.

5.2.1. Idle

In an Idle state, vishing applications exploit Duplicated Contacts with procedures described in Figure 4. The attack begins while the device is in an IDLE state (i.e., PHONE_STATE_IDLE) with no incoming calls ①. The attacker searches for names in the contact list that match predefined target names (e.g., mother, mom, mum, mama, etc.) ②. If a target name is found, the attacker either saves their own phone number under the same name as the target contact or adds their number to the existing target contact ③. Consequently, when an incoming call event occurs, the attacker's number will appear under the target name. And if the user is not vigilant, the attack can go unnoticed. Additionally, when saving a number with an identical name, the Android system does not display a notification or message, potentially preventing the user from realizing that an unwanted contact has been added.




Figure 4. Duplicated Contacts Attack operation.

5.2.2. Incoming Call

A vishing application leverages Display Overlay to alter the caller ID displayed when an attacker places a call to the user. This attack follows the steps outlined in Figure 5. Initially, the application monitors changes in phone state to detect when an incoming call event begins ①. Once an incoming call event triggers and the phone state changes to

`EXTRA_STATE_RINGING`, the attacker checks if the incoming call number matches the pre-set number using `EXTRA_PHONE_NUMBER` ②. If the incoming number aligns with the attacker's preset, a fake incoming call screen, embedded within the application, overlays the display, thereby executing the Display Overlay Attack.




Figure 5. Display Overlay Attack operation.

5.2.3. Outgoing Call

Finally, if the user attempts to dial a number targeted by the attacker, Call Redirection is employed. The vishing application conducts this attack step-by-step, as described in Figure 6. First, when the user presses the dial button, the vishing application checks if the outgoing call event has started using `getAction` ①. If the call state is `ACTION_NEW_OUTGOING_CALL`, the application retrieves the number stored in `EXTRA_PHONE_NUMBER` ②. Should this number match the attacker's target, the vishing application uses `setResultData()` to replace the dialed number with one designated by the attacker ③. Once the outgoing number is altered, the number displayed to the user also changes, necessitating the use of Display Overlay to re-display the originally dialed number. This re-display occurs when the phone state switches to `EXTRA_STATE_OFFHOOK`, overlaying a fake incoming call screen to mask the redirection, thereby preventing the user from detecting any manipulation.




Figure 6. Call Redirection Attack operation.

As a result, when we run these benchmarks alongside VENTINEL, we observe that normal defense operations are maintained, demonstrating that VENTINEL is effective at defending APIs used at API level 28 and below, as well as those at API level 29 and above.

5.3. Comparative Analysis

We conducted a comparative evaluation of vishing defense applications registered on the Google Play Store against VENTINEL using the benchmarking applications we developed, as outlined in Section 5.2. Additionally, to prevent false positives from approved applications exhibiting behavior similar to Display Overlay Attacks, we analyzed the top four most downloaded call screen theme applications to assess whether they were detected as vishing applications. For evaluation, we categorized detection outcomes as follows: detecting a approved application as benign was classified as a True Negative (TN), detecting a approved application as malicious behavior was classified as a False Negative (FN), detecting a benchmark application as malicious behavior was classified as a True Positive

(TP), and detecting a benchmark application as benign behavior was classified as a False Positive (FP). As shown in Table 4, most commercial applications available on the Google Play Store failed to detect the malicious behaviors employed by vishing applications. Some applications that identify these malicious behaviors provide warnings by displaying the incoming or outgoing phone number to the user through pop-up windows when call events occur. Additionally, certain applications compare the incoming phone number at the start of a call with the number displayed when the call transitions to the OFFHOOK state, alerting the user to any discrepancies. However, these approaches are generally limited, as they primarily provide indirect warnings to users or are confined to detecting specific malicious behaviors only. Consequently, to compare defensive measures and address their shortcomings, we manually analyzed the Safe Voice and Call Blocker applications.

Table 4. Detection rate comparison across various applications, including VENTINEL.

	Call Screen Thema Apps				Call Redirection		Display Overlay	Duplicated Contacts
	App 1	App 2	App 3	App 4	API 28 Down	API 29 Up		
WhoWho [11]	TN	TN	TN	TN	TP	TP	TP	FP
PhishingEyes [16]	TN	TN	TN	TN	FP	FP	FP	FP
SmartAntiPhishing [19]	TN	TN	TN	TN	FP	FP	FP	FP
CallApp [17]	TN	TN	TN	TN	FP	FP	FP	FP
Whoscall [15]	TN	TN	TN	TN	TP	TP	TP	FP
Truecaller [12]	TN	TN	TN	TN	TP	TP	TP	FP
SafeVoice [13]	TN	TN	TN	TN	FP, TP	FP, TP	TP	FP
Call Blocker [36]	TN	TN	TN	TN	TP	TP	TP	FP
Show Caller ID & Spam Blocker [18]	TN	TN	TN	TN	FP	FP	FP	FP
Ventinel	TN	TN	TN	TN	TP	TP	TP	TP

TP: True Positive, TN: True Negative, FP: False Positive, FN: False Negative.

5.3.1. SafeVoice

The primary goal of the SafeVoice application is to verify whether the phone number has been altered during outgoing and incoming events and to notify the user accordingly [13]. The application first checks whether phone-related permissions are granted in Call Modulation and monitors for outgoing events. If an ACTION_NEW_OUTGOING_CALL event occurs, the app retrieves the phone number using EXTRA_PHONE_NUMBER, attaches a specific tag to it, and stores it in Shared Preferences. This stored phone number is then saved in the Result using getSharedPreference(tag), while the outgoing number is stored in CallModulation using getStringExtra. Finally, if a discrepancy is detected between the Result (i.e., Original Number) and the CallModulation (i.e., Outgoing Number), the application alerts the user through alert. The operational flow observed in the code presented in Listing 1.

At the time of an outgoing call, SafeVoice displays EXTRA_PHONE_NUMBER alongside getStringExtra(IncomingNumber) to the user. However, this process fails to compare the number collected during the outgoing event with any other numbers before displaying it to the user, limiting its ability to detect Display Overlay Attacks that occur during incoming events. Additionally, because the values stored in Shared Preferences do not change until the next outgoing call occurs, any subsequent legitimate incoming call incorrectly be detected as tampered.

Listing 1. SafeVoice detection code.

```

1 public void onReceive(Context context Intent intent) {
2     if (intent.getAction() != null) {
3         oriNum = "origPhoneNum";
4         ...
5         if (intent.getAction().equals
6             ("android.intent.action.NEW_OUTGOING_CALL"))
7             {this.prefs.edit().putString(oriNum, intent.getStringExtra
8             ("android.intent.extra.PHONE_NUMBER")).apply();
9         }
10        ...
11    }
12    inNum = intent.getStringExtra(inNum);
13    Result = "outgoing changed" + formatNum(Result) + "to" + formatNum(inNum);
14    SafeVoiceReceiver.number = inNum;
15    this.alert(Result, "Call redirection");
16 }
17

```

5.3.2. Call Blocker

The Call Blocker application aims to display the appropriate phone number to the user via a pop-up based on the current phone state [36]. To achieve this, the application checks whether the phone state (i.e., CurrentState) is RINGING or OFFHOOK using getStringExtra(state). It then retrieves the phone number with getStringExtra(IncomingNumber) and presents it to the user. This process can be observed in the code provided in Listing 2. Users verify the phone number displayed by Call Blocker in the pop-up when either an incoming or outgoing call event occurs. However, the application fails to detect any alterations or provide warnings when a manipulation takes place, making it incapable of directly identifying Display Overlay and Call Redirection Attacks. Thus, while Call Blocker effectively shows the relevant phone number during call events, its inability to alert users to potential tampering or manipulations limits its effectiveness as a vishing defense tool.

Listing 2. Call Blocker detection code.

```

1     public void onReceive(Context context Intent intent) {
2         ...
3         if (intent.getAction().equals("PHONE_STATE")) {
4             ...
5             String lastnum = DB.getString(lastnum, "");
6             String curNum = intent.getStringExtra("state");
7             ...
8             if (!TelephonyManager.EXTRA_STATE_RINGING.equals(curNum)) {
9                 if (TelephonyManager.EXTRA_STATE_OFFHOOK.equals(curNum)) {
10                     ...
11                 }
12                 String inNum = intent.getStringExtra("incoming_number");
13             }
14         }
15     }
16

```

5.4. User Study

To validate the proposed technique, we recruited 100 participants for Group A, who received vishing applications, and another 100 participants for Group B, who were provided with both vishing applications and VENTINEL. All participants used Android devices, with the Android version randomly selected. Participants were recruited under the pretext

of a study unrelated to vishing, ensuring they would not recognize the functioning of the vishing applications. After the experiment concluded, we informed them about the study's true nature and requested their participation in a survey regarding vishing. The experiment lasted one week, during which both incoming and outgoing calls were executed randomly. An incoming call was considered unanswered if the participant did not respond, while an outgoing call was counted as unanswered if the subject ended the call within five seconds after answering. As shown in Figure 7, Group A had an average response rate of approximately 64.7% across ten tests, whereas Group B exhibited an average response rate of only 8.9%. Each corresponding number represents the total number of responses in Group A and Group B divided by the number of trials. These results demonstrate that VENTINEL effectively defends against vishing applications.




Figure 7. Response rate comparison between Group A (vishing app) and Group B (vishing app + VENTINEL).

6. Discussion

VENTINEL is engineered to detect attacks from vishing applications (Call Redirection, Display Overlay, and Duplicated Contacts) while operating in the background without necessitating modifications to the Android OS. Vishing applications utilize diverse attack strategies across different Android API levels. To address this, we conducted a detailed analysis of the malicious behavior patterns associated with various API levels and developed VENTINEL to effectively identify and mitigate these threats.

6.1. Lacking Visibility

In Section 5.4, Group B, which used VENTINEL, exhibited an average response rate of 8.9%. Although the overall figure is notably low, we analyzed the experiment to understand why a specific part of Group B exhibited the high response rate. The analysis indicated that the signaller's warning window lacked sufficient visibility. Additionally, the widespread use of Bluetooth earphones and wearable devices led to instances where participants failed to notice the warnings. To address this issue, it is essential to enhance the delivery of alerts related to malicious activities by incorporating voice prompts or establishing a connection between the wearable device and the user's mobile phone for more effective notifications.

6.2. Optimize Performance

We used ML Kit, which requires API level 21 or higher. If your device runs an API level below 21, this feature will not be available. Devices with API levels below 21 need

to rely on alternative solutions like OpenCV. Additionally, since ML Kit's performance has been tested on a limited range of devices, it may encounter issues on less powerful hardware. Therefore, optimization is necessary for these devices.

6.3. APIs Permission Settings

VENTINEL functions entirely within the confines of the application and requires explicit user authorization to operate effectively. It is essential that we provide appropriate notification to users to ensure that all necessary permissions are granted. Furthermore, VENTINEL utilizes APIs that Google has mandated be disabled starting from API level 29. Consequently, if the minimum API level is set to 29 or higher, this feature will be completely deactivated. To address this limitation, it is essential to switch to an alternative API that gives the same results.

6.4. Relay Station Attack

Ventinel is designed to protect against the malicious techniques employed by vishing applications. However, it is important to note that if a relay station alters the phone number, VENTINEL will be unable to detect this modification. To mitigate this limitation, it is necessary to utilize the relay station detection APIs provided by services such as Twilio [37], Hiya [38], YouMail [39], and others. Since domestic services are not supported, close collaboration with telecommunications providers such as SK Telecom [40], LG Telecom [41], and KT Telecom [42] will be essential for effective implementation.

7. Conclusions

This study analyzes the APIs used by vishing applications and reveals that malicious behaviors, such as Call Redirection, Display Overlay and Duplicated Contacts are still be executed using specific APIs available at API level 29 and above. This finding confirms the vulnerabilities associated with APIs modifications and additions presented in the HearMeOut [4]. To address these issues, we propose VENTINEL. We evaluated the detection accuracy of VENTINEL using benchmarks from API level 28 and below, as well as those from API level 29 and above, achieving a high detection rate. Additionally, the application is designed to operate with minimal permissions at the app level, making it easy for users to install and use. Finally, user studies demonstrate that VENTINEL is a robust tool capable of accurately detecting and warning against vishing applications. This research contributes to a better understanding of vishing and assists future studies in this domain.

Author Contributions: Conceptualization, D.K.; methodology, D.K.; software, D.K. and S.O.; validation, D.K.; formal analysis, D.K.; investigation, D.K.; resources, D.K. and S.O.; data curation, D.K.; writing—original draft preparation, D.K., S.O., Y.B., K.J. and H.C.; supervision, J.P., K.J. and H.C.; project administration, H.C. All authors have read and agreed to the published version of the manuscript.

Funding: This work was supported by Institute of Information & Communications Technology Planning & Evaluation (IITP) grant funded by the Korean government (MSIT) (No. RS-2024-00398353, Development of Countermeasure Technologies for Generative AI Security Threats).

Data Availability Statement: Data are contained within the article.

Conflicts of Interest: The authors declare no conflicts of interest.

References

1. Kolandaishamy, R.; Rajagopal, H.; Kolandaishamy, I. Emergence of Cybercrimes in Online Social Networks. In Proceedings of the International Conference on Artificial Life and Robotics, Singapore, 5–7 July 2024; ALife Robotics Co., Ltd.: Oita, Japan, 2024.

2. The Most Common Types of Cyber Crime. Available online : <https://www.statista.com/chart/24593/most-common-types-of-cyber-crime/> (accessed on 25 April 2024).
3. Hashmi, S.I.; George, N.; Saqib, E.; Ali, F.; Siddique, N.; Kashif, S.; Ali, S.; Bajwa, N.U.H.; Javed, M. Training Users to Recognize Persuasion Techniques in Vishing Calls. In Proceedings of the Extended Abstracts of the 2023 CHI Conference on Human Factors in Computing Systems, Hamburg Germany, 23–28 April 2023; pp. 1–8.
4. Kim, J.; Kim, J.; Wi, S.; Kim, Y.; Son, S. HearMeOut: Detecting voice phishing activities in Android. In Proceedings of the 20th Annual International Conference on Mobile Systems, Applications and Services, New York, NY, USA, 27 June–1 July 2022; pp. 422–435.
5. Das, S.; Ganguly, D. Protecting Your Assets: Effective Use of Cybersecurity Measures in Banking Industries. In *Next-Generation Cybersecurity: AI, ML, and Blockchain*; Springer: Berlin/Heidelberg, Germany, 2024; pp. 265–286.
6. Tu, H.; Doupé, A.; Zhao, Z.; Ahn, G.J. Users really do answer telephone scams. In Proceedings of the 28th USENIX Security Symposium (USENIX Security 19), Santa Clara, CA, USA, 14–16 August 2019; pp. 1327–1340.
7. Yeboah-Boateng, E.O.; Amanor, P.M. Phishing, SMiShing & Vishing: An assessment of threats against mobile devices. *J. Emerg. Trends Comput. Inf. Sci.* **2014**, *5*, 297–307.
8. Maggi, F. Are the con artists back? A preliminary analysis of modern phone frauds. In Proceedings of the 2010 10th IEEE International Conference on Computer and Information Technology, Bradford, UK, 29 June–1 July 2010; pp. 824–831.
9. Figueiredo, J.; Carvalho, A.; Castro, D.; Gonçalves, D.; Santos, N. On the Feasibility of Fully AI-automated Vishing Attacks. *arXiv* **2024**, arXiv:2409.13793.
10. Chin, A.G.; Etudo, U.; Harris, M.A. On mobile device security practices and training efficacy: An empirical study. *Inform. Educ.* **2016**, *15*, 235. [[CrossRef](#)]
11. Whowho. Available online: <https://play.google.com/store/apps/details?id=com.ktcs.whowho&hl> (accessed on 7 October 2024).
12. Truecaller. Available online: <https://play.google.com/store/apps/details?id=com.truecaller&hl> (accessed 7 October 2024).
13. SafeVoice. Available online: <https://play.google.com/store/apps/details?id=co.safevoice&hl> (accessed on 20 August 2024).
14. Hiya. Available online: <https://play.google.com/store/apps/details?id=com.webascender.callerid&hl> (accessed on 7 October 2024).
15. Whoscall. Available online: <https://play.google.com/store/apps/details?id=gogolook.callgogolook2> (accessed on 7 October 2024).
16. Phishingeyes. Available online: <https://play.google.com/store/apps/details?id=com.infinigru.lite.phishingeyes&hl> (accessed on 7 October 2024).
17. CallApp. Available online: <https://play.google.com/store/apps/details?id=com.callapp.contacts> (accessed on 7 October 2024).
18. Show Caller ID & Spam Blocker. Available online: <https://play.google.com/store/apps/details?id=com.truecaller.callerid.callername&hl> (accessed on 7 October 2024).
19. SmartAntiPhishing. Available online: <https://play.google.com/store/apps/details?id=com.infinigru.sap.phishingeyes&hl> (accessed on 7 October 2024).
20. Khonji, M.; Iraqi, Y.; Jones, A. Phishing detection: A literature survey. *IEEE Commun. Surv. Tutorials* **2013**, *15*, 2091–2121. [[CrossRef](#)]
21. Naqvi, B.; Perova, K.; Farooq, A.; Makhdoom, I.; Oyedele, S.; Porras, J. Mitigation strategies against the phishing attacks: A systematic literature review. *Comput. Secur.* **2023**, *132*, 103387. [[CrossRef](#)]
22. Wei, X.; Gomez, L.; Neamtiu, I.; Faloutsos, M. Malicious android applications in the enterprise: What do they do and how do we fix it? In Proceedings of the 2012 IEEE 28th International Conference on Data Engineering Workshops, Arlington, VA, USA, 1–5 April 2012; pp. 251–254.
23. Android Developer PopupWindow. Available online: <https://developer.android.com/reference/android/widget/PopupWindow> (accessed on 7 October 2024).
24. Android Developer WindowManager. Available online: <https://developer.android.com/reference/android/view/WindowManager> (accessed on 7 October 2024).
25. Finley, J.R.; Naaz, F.; Goh, F.W.; Finley, J.R.; Naaz, F.; Goh, F.W. Results: Behaviors and experiences with internal and external memory. In *Memory and Technology: How We Use Information in the Brain and the World*; Springer: Berlin/Heidelberg, Germany, 2018; pp. 25–48.
26. Muhammad, Z.; Anwar, Z.; Javed, A.R.; Saleem, B.; Abbas, S.; Gadekallu, T.R. Smartphone Security and Privacy: A Survey on APTs, Sensor-Based Attacks, Side-Channel Attacks, Google Play Attacks, and Defenses. *Technologies* **2023**, *11*, 76. [[CrossRef](#)]
27. Hutchinson, S.; Zhou, B.; Karabiyik, U. Are we really protected? An investigation into the play protect service. In Proceedings of the 2019 IEEE International Conference on Big Data (Big Data), Los Angeles, CA, USA, 9–12 December 2019; pp. 4997–5004.
28. Griffin, S.E.; Rackley, C.C. Vishing. In Proceedings of the 5th Annual Conference on Information Security Curriculum Development, Kennesaw, Georgia, 26–27 September 2008; pp. 33–35.

29. Jain, A.K.; Gupta, B. A survey of phishing attack techniques, defence mechanisms and open research challenges. *Enterp. Inf. Syst.* **2022**, *16*, 527–565. [[CrossRef](#)]
30. Jones, K.S.; Armstrong, M.E.; Tornblad, M.K.; Namin, A.S. How social engineers use persuasion principles during vishing attacks. *Inf. Comput. Secur.* **2020**, *29*, 314–331. [[CrossRef](#)]
31. Phishing Activity Trends Report 4th Quarter 2023. Available online: <https://apwg.org/trendsreports/> (accessed on 25 April 2024).
32. Android Security Bulletins. Available online: <https://source.android.com/docs/security/bulletin/asb-overview?hl> (accessed on 7 October 2024).
33. All Android Releases. Available online: <https://developer.android.com/about/versions> (accessed on 7 October 2024).
34. Android Developers Android 9 Changes. Available online: <https://developer.android.com/about/versions/pie/android-9.0-changes-all> (accessed on 7 October 2024).
35. Google ML Kit . Available online: <https://developers.google.com/ml-kit> (accessed on 12 June 2024).
36. Call Blocker. Available online: <https://play.google.com/store/apps/details?id=call.blacklist.blocker&hl> (accessed on 20 August 2024).
37. Twilio. Available online: <https://www.twilio.com/en-us> (accessed on 1 November 2024).
38. Hiya. Available online: <https://www.hiya.com/> (accessed on 1 November 2024).
39. YouMail. Available online: <https://www.youmail.com/> (accessed on 1 November 2024).
40. SK Telecom. Available online: https://www.sktelecom.com/index_en.html (accessed on 1 November 2024).
41. LG Telecom. Available online: <https://www.lguplus.com/> (accessed on 1 November 2024).
42. KT Telecom. Available online: <https://www.kt.com/> (accessed on 1 November 2024).

Disclaimer/Publisher's Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.